

Q: "Is the adult in the image chasing a child?" (ground truth: Yes)

ringed cells decode to a word from the question (adult, chasing, child); Qwen3-VL-8B, layer 32


Input



every image patch is shown at right


SIT (question-last)

→ "Yes" ✓



STI (question-first)

→ "No" ✗



STIT (ours, echoing)

→ "Yes" ✓

